

## Table of Contents

1. Count(*).....	1
2. Count(expression) .....	1
3. Count(Distinct expression) .....	2
4. Count with Inner and with Outer Joins .....	5
5. Summary .....	6

I am putting this discussion of Count in a separate document because Count has some useful features different than Max, Min, Sum, and Avg.

## 1. Count(\*)

The expression count(\*) will return the number of rows in the table expression. If the table is empty- has no rows- count(\*) will return 0. If there are rows, count(\*) will tell you how many rows are in the table.

Demo 01: Number of rows in the customer table

```
select COUNT(*) As "Number of customers"
from Customer.customers;
+-----+
| Number of customers |
+-----+
|                   34 |
+-----+
```

## 2. Count(expression)

Count(expression) will determine the number of rows in the table expression where that expression is not null.

Demo 02: Using the Count aggregate function with a column attribute.

At the time this query was run, we had 34 customers- all customers have a customer\_id value since that is the primary key for this table; 33 of those customers have a value for first name.

```
select
  COUNT(customer_id) As "Number of customers"
, COUNT(customer_name_first) As "Number with a first name"
from Customer.customers;
+-----+-----+
| Number of customers | Number with a first name |
+-----+-----+
|                   34 |                   33 |
+-----+-----+
```

If the table is empty, or if all of the rows have a null for the expression, count(expression) will return 0. Remember that Max, Min, Sum, and Avg returned a null in that case. If the table expression is empty. we know how many rows it has- 0.

Demo 03: Currently the largest salary value in the employee table is 120000. If I run the following query filtering for salary values equal to 200000. I get an empty set

```
select salary
from Employee.employees
where salary = 200000;
Empty set
```

But if I do the count and sum of those salaries I get a count of 0 and a sum of null.

```
select sum(salary) as SumSalary, count(salary) as CountSalary
from Employee.employees
where salary = 200000;
+-----+-----+
| SumSalary | CountSalary |
+-----+-----+
|          | 0           |
+-----+-----+
1 row in set (0.00 sec)
```

Some people think this does not make a lot of sense, but that is the way these functions work.

You will sometimes hear people suggest that you use an expression such as `count(1)` to count rows, or `count(customer_id)` with the customer table to count rows - `customer_id` is the pk and is never null. Those will work but there is no advantage to doing this. You could count any non null column- such as `Count(customer_name_last)` but that is confusing and someone might have changed that column to be a nullable column without checking all queries that use that table. Use a sensible column for count- if you are supposed to count customers, use `count(customer_id)`; if you are supposed to count products, use `count(prod_id)`; if you are supposed to count employees, use `count(emp_id)`. But also see the next section of `count(distinct...)`

`Count(1)` works because 1 is a numeric literal- but you could use `count(0)` or `count(3.14159)` or `count('Hi There')` the same way since they are all literals. (If you are working at a job and your boss says to use `count(1)` instead of `count(*)`, or if that is a company standard- go ahead. It won't hurt the query result.)

Demo 04: We can use count with an expression. Suppose we wanted to know how many customers we have and how many have a last name starting with the letters 'Mor'

```
select Count(customer_name_last) as CustCount
, count(case when customer_name_last like 'Mor%' then 1 end) as CustNameMorX
from Customer.customers;
+-----+-----+
| CustCount | CustNameMorX |
+-----+-----+
|          | 5            |
+-----+-----+
1 row in set (0.03 sec)
```

### 3. Count(Distinct expression)

Suppose we want to know how many products we have currently on order. It is important to realize that a simple request like this could have different interpretations. Suppose we have only these rows in our order details table.

Ord_id	Line_Item_id	Prod_id	Quantity	Quoted_price
1	1	P12	2	5.00
1	2	T20	4	12.00
1	3	R67	1	25.00
2	1	P12	2	5.00
3	1	P12	8	6.50
3	2	R67	5	25.00

The question "how many products do we have currently on order ?" could be interpreted to mean "how many line items do we have ?" ; in this example we have 6 line items- we could answer this by using Count(\*). The question could also be interpreted to mean "what is the total quantity of items that we have on order?"; this would be  $2 + 4 + 1 + 2 + 8 + 5 = 22$  – for that use the function Sum (quantity). Another interpretation is "how many different products do we have on order "; this would be products P12, R67, and T20; we have three different products currently on order. This is answered using the Count (Distinct prod\_id) function.

Demo 05: Using count distinct . We have 184 rows in the order headers table and 34 different products have been ordered.

```
select count(distinct prod_id) as CountDistinctProducts
, count(*) as CountProducts
from OrderEntry.orderDetails;
+-----+-----+
| CountDistinctProducts | CountProducts |
+-----+-----+
| 34 | 184 |
+-----+-----+
```

Note that the use of this keyword Distinct is different in syntax and meaning than the use of Distinct with a Select statement to return only unique rows.

When you add Distinct **inside** the argument list to an **aggregate** function, the function will determine the answer by using only unique values.

Demo 06: How many orders do we have? We can count the order ids in the order headers table.

```
select count( order_id)
from OrderEntry.orderHeaders;
+-----+
| count( order_id) |
+-----+
| 97 |
+-----+
```

Demo 07: If we count the order id in the order detail tables, we count any order with multiple order lines more than once. Use distinct to count each order only once.

```
select count(order_id), count(distinct order_id)
from OrderEntry.orderDetails;
+-----+-----+
| count(order_id) | count(distinct order_id) |
+-----+-----+
| 184 | 94 |
+-----+-----+
```

Demo 08: How many customers do we have with orders? We can count the distinct customer ids in the order headers table. If we count customer\_id, then we count customer with multiple orders multiple times. We have only 34 customers.

```
select count(distinct customer_id), count(customer_id)
from OrderEntry.orderHeaders;
+-----+-----+
| count(distinct customer_id) | count(customer_id) |
+-----+-----+
| 21 | 97 |
+-----+-----+
```

Demo 09: (MySQL feature) Suppose we want to find out how many different combinations of customers and shipping modes we have. I am not going to count the orders where there is no info on the shipping modes.

```
select count( distinct customer_id, shipping_mode_id)
from OrderEntry.orderHeaders
where shipping_mode_id is not null;
+-----+
| count( distinct customer_id, shipping_mode_id) |
+-----+
|                                     33 |
+-----+
```

Demo 10: Alternate approaches that work in more dbms.

```
select  COUNT(*)
from (
    select distinct customer_id, shipping_mode_id
    from OrderEntry.orderHeaders
    where shipping_mode_id is not null) dataSource;
+-----+
| COUNT(*) |
+-----+
|         33 |
+-----+
```

Demo 11: We have about 100 rows in the order headers table. We want to see how many different months are included in those orders.

```
select count( distinct concat(year(order_date), '-',month(order_date))) as
"number of months"
from OrderEntry.orderHeaders
where year(order_date) = 2015;
+-----+
| number of months |
+-----+
|                 7 |
+-----+
```

Demo 12: I put the data gathering query is a subquery in the form clause. ( we have not discussed this yet) That lets the main query deal with the aggregates. Be careful to choose a format for dates that have a two digit month value. These are strings and you need the leading zero so that Feb 02 comes before Nov 11.

```
select count(YrMnth) as "number of months", min(YrMnth) as "earliest month"
, max(YrMnth) as "latest month"
from (
    select distinct Date_format(order_date, '%Y-%m') as YrMnth
    from OrderEntry.orderHeaders
    where year(order_date) = 2015
    ) dataSource;
+-----+-----+-----+
| number of months | earliest month | latest month |
+-----+-----+-----+
|                 7 | 2015-04       | 2015-12     |
+-----+-----+-----+
1 row in set (0.00 sec)
```

You can use Distinct with any of the aggregates- but don't use it if you really don't need it. If you are looking for the price of the most expensive item we sell, Max(price) will return the same value as Max(Distinct price).

I have never seen a good business reason for calculating sum(distinct prod\_list\_price) or for calculating sum(prod\_list\_price) from the products table.

## 4. Count with Inner and with Outer Joins

Suppose we join the tables products and order details to see which products have been ordered. If we do this with an inner join then we get only products which have been ordered and if we use a left join (Products Left Join order\_details) then we get products with orders and products without orders. This means you need to take more care with aggregate functions. The Count function is the easiest function for illustrating this point.

### Demo 13: Inner join

```
select PR.prod_id, OD.prod_id, PR.catg_id
from Product.products PR
left join OrderEntry.orderDetails OD on Pr.prod_id = OD.prod_id
where pr.prod_id between 1140 and 1150
order by PR.prod_id;
```

prod_id	prod_id	catg_id
1140	1140	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1150	1150	PET
1150	1150	PET
1150	1150	PET
1150	1150	PET
1150	1150	PET

16 rows in set (0.00 sec)

### Demo 14: Outer join- we get 2 more rows - products 1142, 1143 with no orders

```
select PR.prod_id, OD.prod_id, PR.catg_id
from a_prd.products PR
left join OrderEntry.orderDetails OD on Pr.prod_id = OD.prod_id
where pr.prod_id between 1140 and 1150
order by PR.prod_id;
```

prod_id	prod_id	catg_id
1140	1140	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1141	1141	PET
1142	NULL	PET
1143	NULL	PET

```

| 1150 | 1150 | PET |
| 1150 | 1150 | PET |
| 1150 | 1150 | PET |
| 1150 | 1150 | PET |
| 1150 | 1150 | PET |
+-----+-----+-----+
18 rows in set (0.00 sec)

```

#### Demo 15: using the inner join and aggregates

```

select count(*) as "RowCount"
, count(distinct order_id) as OrderCount
, count(distinct PR.prod_id) as "PR_DistProdCount"
, count(distinct OD.prod_id) as "OD_DistProdCount"
from Product.products PR
join OrderEntry.orderDetails OD on PR.prod_id = OD.prod_id
where pr.prod_id between 1140 and 1150;
+-----+-----+-----+-----+
| RowCount | OrderCount | PR_DistProdCount | OD_DistProdCount |
+-----+-----+-----+-----+
| 16 | 11 | 3 | 3 |
+-----+-----+-----+-----+

```

#### Demo 16: using the outer join and aggregates

```

select count(*) as "RowCount"
, count(distinct order_id) as OrderCount
, count(distinct PR.prod_id) as "PR_DistProdCount"
, count(distinct OD.prod_id) as "OD_DistProdCount"
from Product.products PR
left join OrderEntry.orderDetails OD on PR.prod_id = OD.prod_id
where pr.prod_id between 1140 and 1150;
+-----+-----+-----+-----+
| RowCount | OrderCount | PR_DistProdCount | OD_DistProdCount |
+-----+-----+-----+-----+
| 18 | 11 | 5 | 3 |
+-----+-----+-----+-----+

```

The differences between the results for Demo 13 and 14.:

**RowCount** With the outer join we get two more rows because we include the rows for the two products with no orders

**OrderCount** These values are the same, we are counting the order\_id values which occur only in the order details table. Since the rows for the products with no orders have a null in that column, they do not get counted.

**PR\_DistProdCount** With this column we are counting the product id value from the products table. With the outer join we do count the product id for the two products with no orders.

**OD\_DistProdCount.** With this column we are counting the product id value from the order details table. Since the rows for the products with no orders have a null in that column, they do not get counted

## 5. Summary

Count (\*) counts rows.

Count (col) counts the values for that column, and does not count Nulls.

Count will always return a numeric integer answer- if the table is empty or there are no matching rows, it will return 0. You don't need to use coalesce with count.

Sum, Avg, Max and Min will return nulls when there are no matching rows

Sum, Avg, Max and Min ignore nulls when doing their calculations

This is a point in the semester where it makes sense to realize that SQL is like every other computer system. It is built on a system of rules and it implements those rules. The rules it has will not always agree with what you think they should be- but it makes no sense to fight a compiler. You need to know the rules the system uses and decide if those rules make sense for the processing you need to do.